

# SLAM++: Simultaneous Localisation and Mapping at the Level of Objects

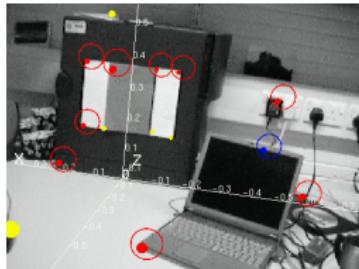
Renato F. Salas-Moreno<sup>1</sup>, Richard A. Newcombe<sup>2</sup>, Hauke Strasdat<sup>1</sup>, Paul H. J. Kelly<sup>1</sup> and Andrew J. Davison<sup>1</sup>

Department of Computing,  
Imperial College London<sup>1</sup>  
&  
CSE, University of Washington<sup>2</sup>

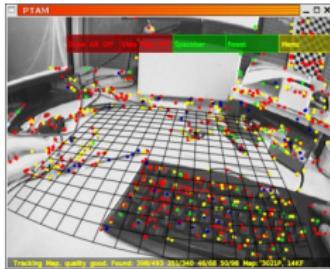
June 28, 2014

- 1 Introduction and Motivation
- 2 SLAM++ System
- 3 Object Level Map Representation
- 4 Real-time 3D Object Pose Detection
- 5 Local Camera Relocalisation

# Real-time Dense SLAM



MonoSLAM, 2003



PTAM, 2007



DTAM, 2011



KinectFusion, 2011

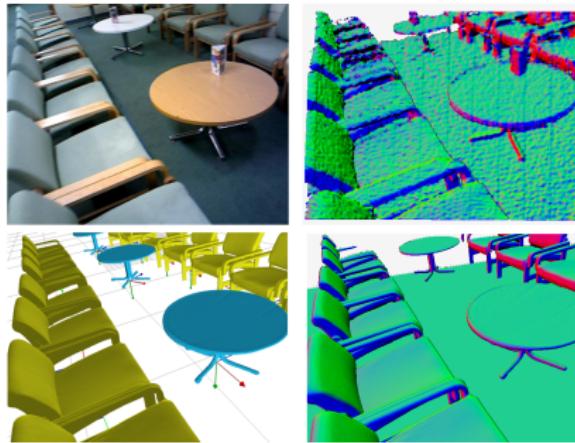
## 3D Representation is Evolving

- Increase in density of representation benefiting from increasing processing and memory resources.
- But also increasingly taking account of real-word priors that the world is not a point cloud.

## To recap

- Maps initially represented by keypoints.
- Density increased towards full dense surfaces.
- World still meaningless: e.g. Keypoint or surface patch on a keyboard same as in monitor as far as interaction is concerned due to limited perception.

# SLAM++: A Pure Object-Level SLAM System



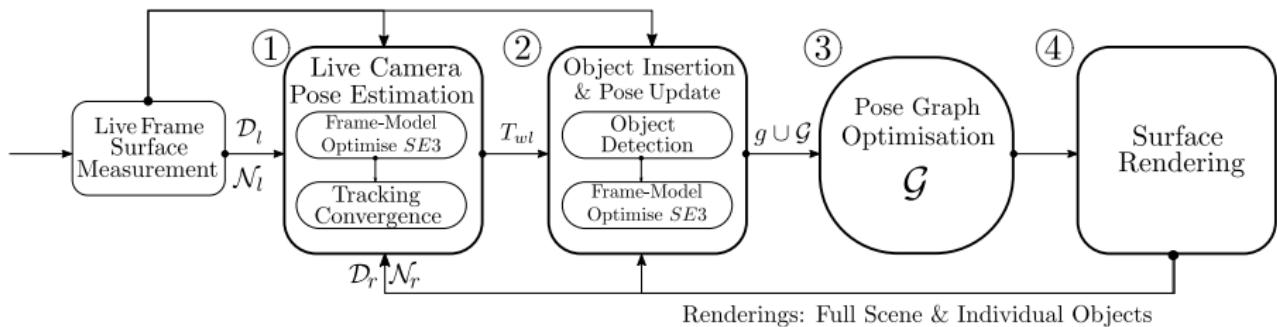
CVPR 2013.

- Brings semantic understanding in the loop of SLAM itself.
- SLAM++: SLAM at the level of objects, Salas-Moreno, Newcombe, Strasdat, Kelly, Davison, CVPR 2013.

By recognizing objects in the loop, we have advantages:

- Instant semantic understanding enables interaction (robotics, AR).
- Dense representation and full predictive power of KinectFusion for robust tracking but very memory-efficient and scalable.
- Enables the inclusion of domain-level knowledge to deal with incorrect mapping.

# System Overview



## Real-Time Loop

- Main components of the system. Mostly on GPU.
- Begin with live measurement from sensor.
- Objects are detected and placed into a map.
- Map used for camera tracking (localisation).
- World represented by objects as nodes in a graph.
- Map is rendered and compared against new measurements to update pose.

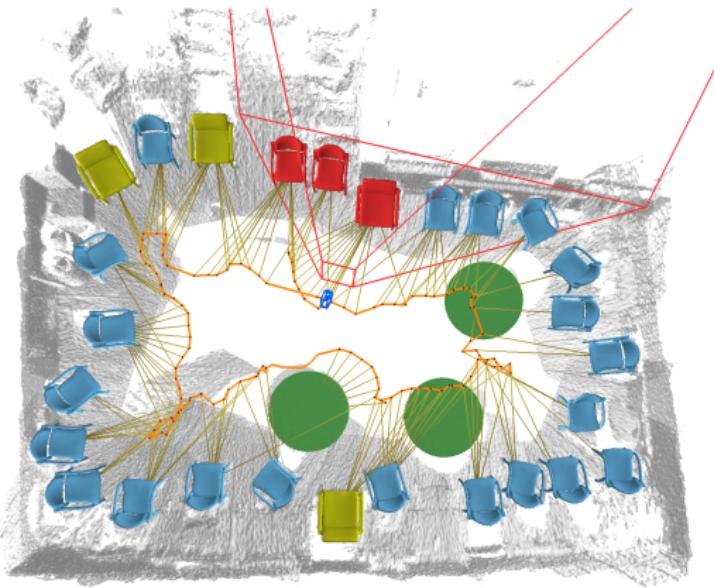
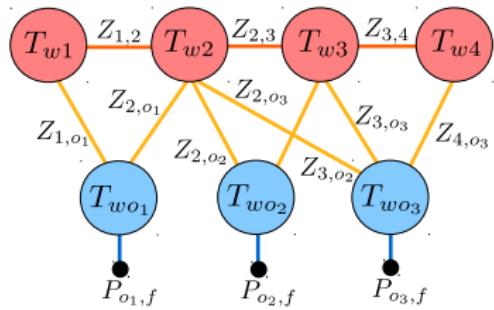
# Object Database



Currently a small number of precisely known objects

- Carefully scanned using KinectFusion and manually segmented.

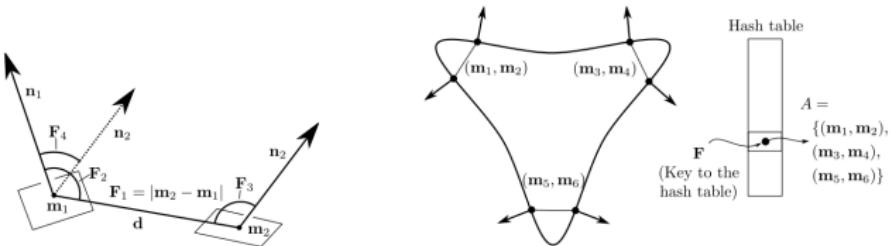
# Map Representation



## A Pure Graph of Objects

- Large room that has been mapped with objects.
- Camera trajectory in orange and measurements in yellow.

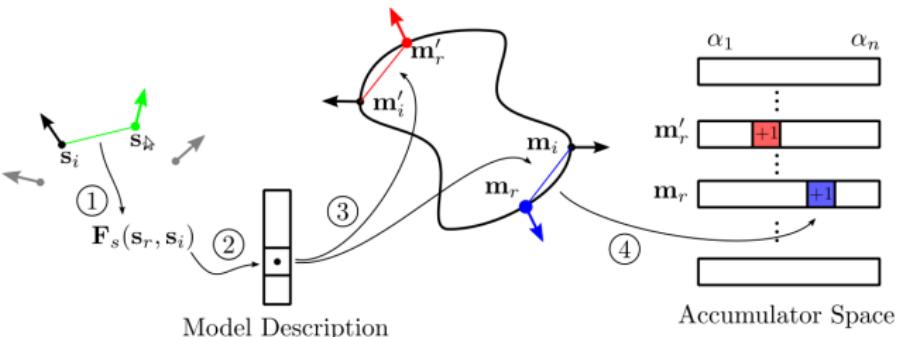
# Real-Time Object Detection



From 'Model Globally, Match Locally', Drost *et al.*, CVPR 2010

- Each object is described by a combination of all point pair features, describing the relative angles and distance between points.
- Object features are discretised and grouped into a Hash Table and used as a lookup for runtime queries coming from depth map features.

# Real-Time Object Detection



From 'Model Globally, Match Locally', Drost *et al.*, CVPR 2010

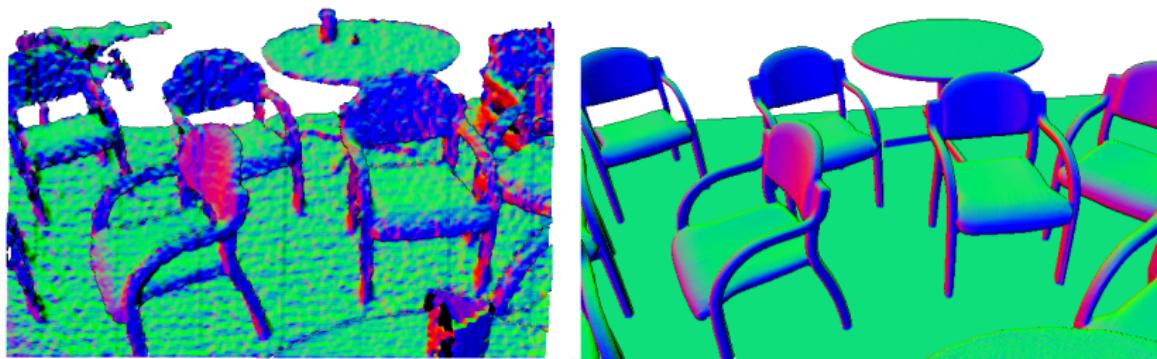
- Feature matches cast votes in accumulator for object pose. The pose with max votes assumed to be a rough estimate of true object pose.
- We have a new efficient parallel implementation on GPU/CPU.
- Still doesn't scale well to a large number of objects.



Noisy raw detected poses refined with ICP against object model

- Detections with the previous step are only approximations due to feature discretisation.
- Pose is refined with Iterative Closest Point that minimise a point plane error metric.

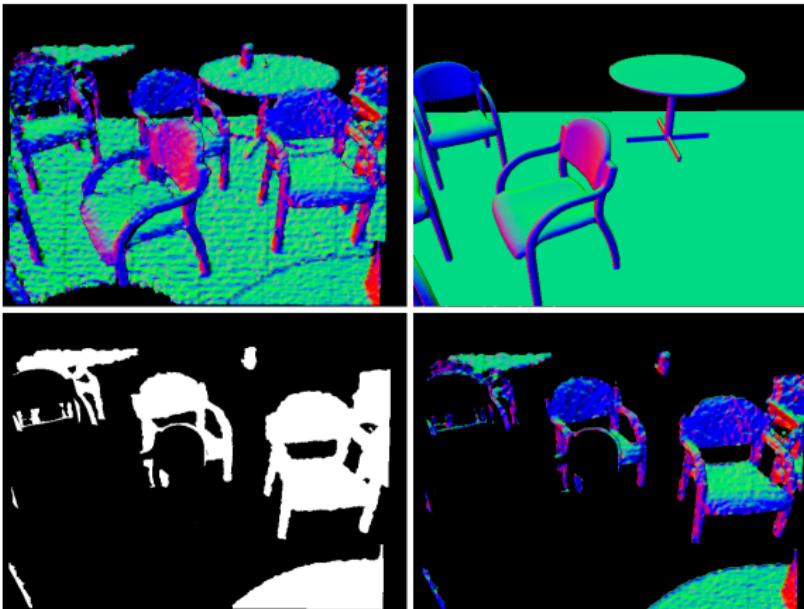
# Camera Tracking



We use ICP again, this time to localise the camera

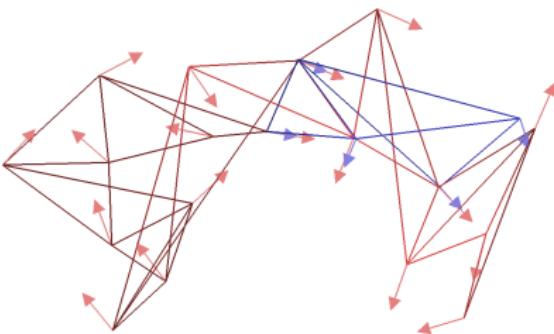
- On the left is the live noisy measurement from the camera.
- On the right is the clean rendered world.
- Assuming the world is static, a misalignment between the Measurement and our Prediction has to be due to camera motion.
- Minimising the alignment error equals to camera pose update.

# Active Object Search in the SLAM Loop



Mask-out already mapped areas

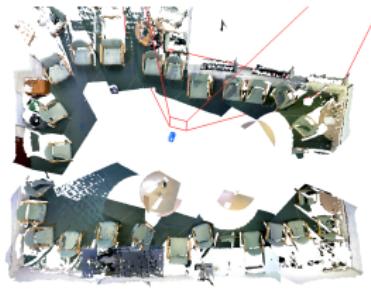
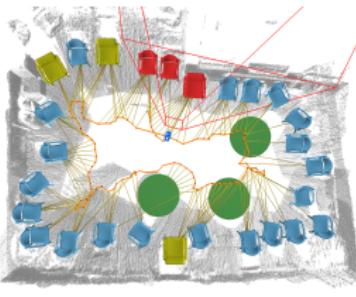
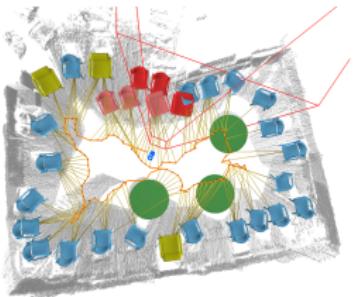
- Detection samples focused on as-yet-unknown regions.
- **Much easier detection** of distant or occluded objects.



## Relocalisation using graph matching

- When the camera is lost, we start mapping again, and when we map at least 3 objects we match the new graph to the main one.
- Each graph is represented as oriented points in a mesh and can be matched with essentially the same code used in object recognition; these graphs are very small typically so this is fast.

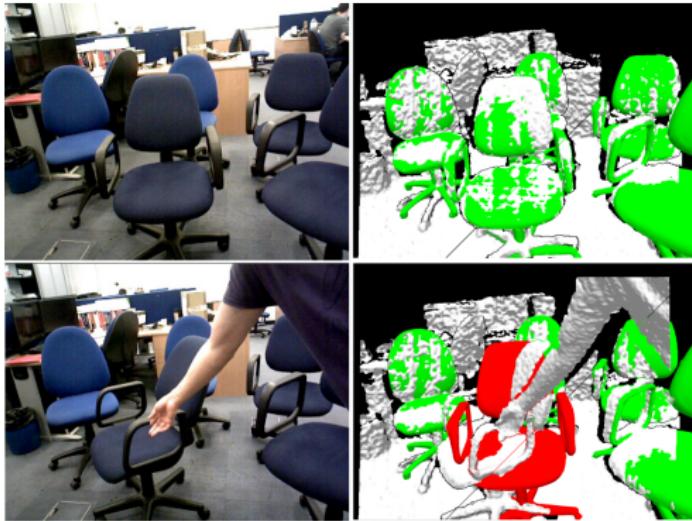
# Large Loops



## Loop closure detection

- When mapping large areas errors accumulate and is difficult to close loops. Set of red chairs not quite on the same place.
- Loops can be detected using the same graph matching mechanism as relocalisation looking for self-similarity in a graph.

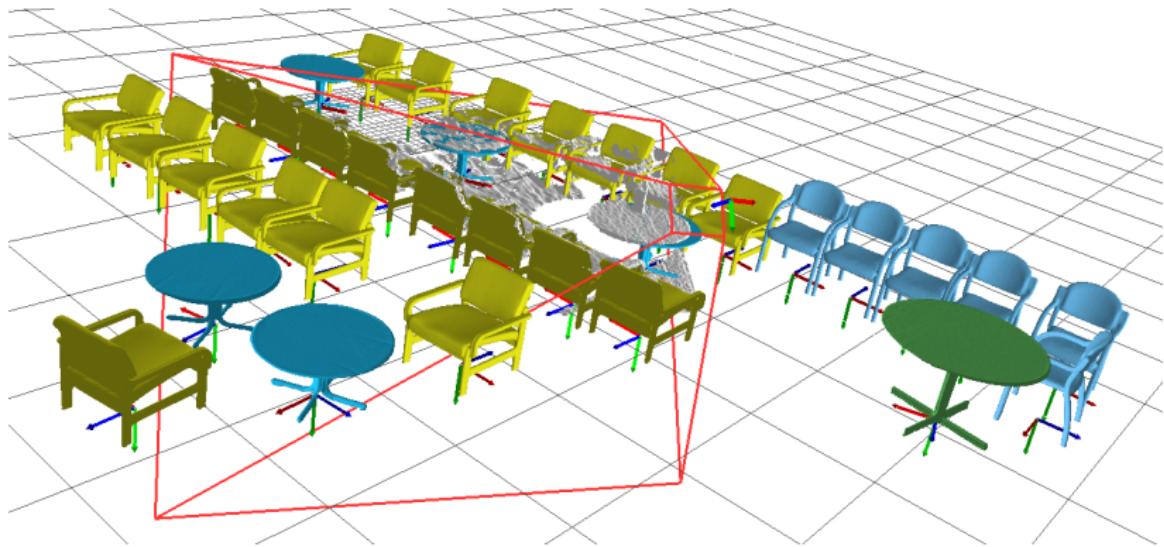
# Moved Object Detection



## Object Moved While Viewed

- We can also easily detect when an object moves using ICP gating on each object individually.

# Whole Room Mapping



## Large Room Rapidly Mapped

- Four object types; object-floor constraints but no object-object layout constraints.

# Augmented Reality



## Context-Aware AR

- Since we know chairs are for sitting and floor for standing we command virtual characters to take a sit or dance.
- Mapped world objects occlude augmentations precisely.
- Other real-world things like people seen in the depth image can also occlude augmentations, but roughly.